Falsification-First
Wiki Editing

A next-generation wiki prototype where every significant claim must separate source, evidence, inference, risk, and — crucially — what would **weaken or falsify the claim**, before it can present itself as credible.

`v0.1 · prototype`
`6 domains · 12 cases`
`deterministic validator`
`AI-structured · human-audited`

THE PROBLEM WITH CITATION-ONLY CREDIBILITY

✗ Current Wikipedia / AI writing

Source → Claim

Citation is necessary, but not sufficient. Source can be real while the claim misrepresents what it says. Inference is invisible. Falsifiability is never asked.

✓ Epistemic Audit Graph

Source → SourceClaim → Evidence → Inference
↓ Hypothesis + Falsifier REQUIRED + Risk +
Narrative

Validators enforce: if a significant hypothesis has no falsification condition, it is **rejected**, not published.

DESIGN PRINCIPLE

AI structures claims

Proposes node decomposition and change requests — does not decide truth

Code enforces audit rules

Deterministic Python validator · domain-specific overclaim rules · no LLM judge

Humans inspect and revise

Editor reviews extracted nodes, falsifier, and risk before accepting the change

Key innovation: The validator is deliberately small and inspectable. The important rules are enforced by Python and JSON rule files — not by asking an LLM whether a claim is true.

VALIDATOR OUTPUT — SAME CLAIM, DIFFERENT STRUCTURE

```
✓ PASS

domain: history / humanities
hypothesis: Jomon male high-status placement possible in some regions
tier: C (testable question)
falsifier: ✓ present
M-tag evidence only: ✓ no
source-claim separation: ✓ separated
overclaim: ✓ none flagged
```

```
✗ REJECT

domain: history / humanities
hypothesis: Jomon conquered and dominated Yayoi society
tier: A (claimed as established)
falsifier: ✗ missing — A-tier requires falsifier
M-tag evidence only: ✗ myth analogy used as direct proof
source-claim separation: ✗ user inference merged into source
overclaim: ✗ domain-generalization flag raised
```

SIX-DOMAIN DEMO — EACH DOMAIN HAS ONE SCOPED (GOOD) CASE AND ONE REJECTED OVERCLAIM CASE

Domain	Scoped (good)	Rejected overclaim
History & Humanities	✓ scoped abductive	✗ myth→certainty
Biomedicine	✓ mouse model	✗ mouse→human
Social Science	✓ association	✗ correlation→cause
Climate & Earth	✓ event+model	✗ one year→trend
AI / Computer Science	✓ benchmark scoped	✗ benchmark→AGI
Law / Policy & Ethics	✓ layers separated	✗ legal→ethical

REPOSITORY

github.com/[username]/epistemic-audit-codex-starter

Static demo: `frontend/index.html`
Validator: `scripts/validate_cases.py`
Six-domain cases: `data/sample_cases/`

LIMITATIONS (HONEST)

- Sample citations are placeholders — need expert review
- No DOI/PDF extraction yet
- No production server, auth, or multi-user editing
- Domain rules are hackathon validation rules, not expert consensus
- The history/humanities case is structurally valid, not endorsed history

ASK

Looking for feedback from WikiCred / Wikimedia credibility builders on:

- Whether **mandatory falsification conditions** fit citation transparency work
- Which existing WikiCred projects to connect to
- Whether this fits a **monthly WikiCred meetup** or Demo Hour